

S&DS 365 / 665
Intermediate Machine Learning

Syncing up in person

Monday, February 7

Yale

Topics for today

- Recap: Everything so far
- Explanations, derivations on board
- Kernel density estimation
- If time, begin neural networks

Notes

- Notes posted to course page
<http://interml.ydata123.org>
- Bias-variance tradeoff for smoothing
- More material on Mercer kernels



Nonparametric Regression

Given $(X_1, Y_1), \dots, (X_n, Y_n)$ predict Y from X .

Assume only that $Y_i = m(X_i) + \epsilon_i$ where $m(x)$ is a smooth function of x .

The most popular methods are *kernel methods*. However, there are two types of kernels:

- 1 Smoothing kernels
- 2 Mercer kernels

Smoothing kernels involve local averaging.
Mercer kernels involve regularization.

Smoothing Kernels

- Smoothing kernel estimator:

$$\hat{m}_h(x) = \frac{\sum_{i=1}^n Y_i K_h(X_i, x)}{\sum_{i=1}^n K_h(X_i, x)}$$

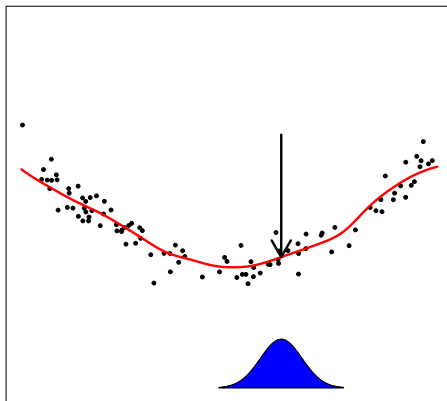
where $K_h(x, z)$ is a *kernel* such as

$$K_h(x, z) = \exp\left(-\frac{\|x - z\|^2}{2h^2}\right)$$


and $h > 0$ is called the *bandwidth*.

- $\hat{m}_h(x)$ is just a local average of the Y_i 's near x .
- The bandwidth h controls the bias-variance tradeoff:
Small h = large variance while *large h = large bias*.


$\hat{m}(x)$ is a local average




Effect of the bandwidth h




very small bandwidth



small bandwidth



medium bandwidth



large bandwidth

Smoothing Kernels

$$\text{Risk} = \mathbb{E}(Y - \hat{m}_h(X))^2 = \text{bias}^2 + \text{variance} + \sigma^2.$$

$$\text{bias}^2 \approx h^4,$$


$$\text{variance} \approx \frac{1}{nh^p} \quad \text{where } p = \text{dimension of } X.$$

$\sigma^2 = \mathbb{E}(Y - m(X))^2$ is the unavoidable prediction error.


small h: low bias, high variance (undersmoothing)

large h: high bias, low variance (oversmoothing)

Risk Versus Bandwidth



h



optimal h

Another Approach: Mercer Kernels

Instead of using local smoothing, we can optimize the fit to the data subject to regularization (penalization). Choose \hat{m} to minimize

$$\sum_i (Y_i - \hat{m}(X_i))^2 + \lambda \text{penalty}(\hat{m})$$

where $\text{penalty}(\hat{m})$ is a *roughness penalty*.

λ is a parameter that controls the amount of smoothing.

How do we construct a penalty that measures roughness? One approach is: *Mercer Kernels* and *RKHS = Reproducing Kernel Hilbert Spaces*.

What is a Mercer Kernel?

A *Mercer Kernel* $K(x, x')$ is symmetric and positive definite:

$$\int \int f(x)f(x')K(x, x') dx dx' \geq 0 \quad \text{for all } f.$$

Example: $K(x, x') = e^{-\|x-x'\|^2/2}$.

Think of $K(x, x')$ as the *similarity* between x and x' . We will create a set of *basis functions* based on K .

Fix z and think of $K(z, x)$ as a function of x . That is,

$$K(z, x) = K_z(x)$$

is a function of the second argument, with the first argument fixed.

Mercer Kernels

Let

$$\mathcal{F} = \left\{ f(\cdot) = \sum_{j=1}^k \beta_j K(z_j, \cdot) \right\}$$

Define a norm: $\|f\|_K = \sum_j \sum_k \beta_j \beta_k K(z_j, z_k)$. $\|f\|_K$ *small means f smooth*.

If $f = \sum_r \alpha_r K(z_r, \cdot)$, $g = \sum_s \beta_s K(w_s, \cdot)$, the inner product is

$$\langle f, g \rangle_K = \sum_r \sum_s \alpha_r \beta_s K(z_r, w_s).$$

\mathcal{F} is a **reproducing kernel Hilbert space (RKHS)** because

$$\langle f, K(x, \cdot) \rangle = f(x)$$

Mercer Kernels

Check that this is well defined:

If $f = \sum_r \alpha_r K(z_r, \cdot)$, $g = \sum_s \beta_s K(w_s, \cdot)$, the inner product is

$$\begin{aligned}\langle f, g \rangle_K &= \sum_r \sum_s \alpha_r \beta_s K(z_r, w_s) \\ &= \sum_r \alpha_r g(z_r) \\ &= \sum_s \beta_s f(w_s)\end{aligned}$$

using the reproducing property $\langle f, K(x, \cdot) \rangle = f(x)$

Nonparametric Regression: Mercer Kernels

Representer Theorem

Let \hat{m} minimize

$$J(m) = \sum_{i=1}^n (Y_i - m(X_i))^2 + \lambda \|m\|_K^2.$$

Then

$$\hat{m}(x) = \sum_{i=1}^n \alpha_i K(X_i, x)$$

for some $\alpha_1, \dots, \alpha_n$.

So, we only need to find the coefficients

$$\alpha = (\alpha_1, \dots, \alpha_n).$$

Nonparametric Regression: Mercer Kernels

In more detail, we can write any $f \in \mathcal{H}_K$ as

$$f = \sum_i \alpha_i K(X_i, \cdot) + v$$

where v is orthogonal to the span of the functions $K(X_i, \cdot)$

By the reproducing property, $f(X_i)$ does not depend on v , and

$$\|f\|_K^2 = \alpha^T \mathbb{K} \alpha + \|v\|_K^2.$$

So, it must be that the minimizing function has $v = 0$.

Nonparametric Regression: Mercer Kernels

Plug $\hat{m}(x) = \sum_{i=1}^n \alpha_i K(X_i, x)$ into J :

$$J(\alpha) = \|Y - \mathbb{K}\alpha\|^2 + \lambda \alpha^T \mathbb{K}\alpha$$

where $\mathbb{K}_{jk} = K(X_j, X_k)$

Now we find α to minimize J . We get: $\hat{\alpha} = (\mathbb{K} + \lambda I)^{-1} Y$ and $\hat{m}(x) = \sum_i \hat{\alpha}_i K(X_i, x)$.

The estimator depends on the amount of regularization λ . Again, there is a bias-variance tradeoff. We choose λ by cross-validation. This is like the bandwidth in smoothing kernel regression.


Smoothing Kernels *Versus* Mercer Kernels

Smoothing kernels: the bandwidth h controls the amount of smoothing.


Mercer kernels: norm $\|f\|_K$ controls the amount of smoothing.

In practice these two methods give answers that are very similar.


Mercer Kernels: Examples




very small lambda



small lambda



medium lambda



large lambda

Multiple Regression

Both methods extend easily to the case where X has dimension $p > 1$. For example, just use

$$K(x, x') = e^{-\|x-x'\|^2/2}.$$

However, this is hard to interpret and is subject to the **curse of dimensionality**. This means that the *statistical performance* and the *computational complexity* degrade as dimension p increases.

An alternative is to use something less nonparametric such as **additive model** where we restrict $m(x_1, \dots, x_p)$ to be of the form:

$$m(x_1, \dots, x_p) = \beta_0 + \sum_j m_j(x_j).$$

Additive Models

Model: $m(x) = \beta_0 + \sum_{j=1}^p m_j(x_j)$.

We can take $\hat{\beta}_0 = \bar{Y}$ and we will ignore β_0 from now on.

We want to minimize

$$\sum_{i=1}^n \left(Y_i - (m_1(X_{i1}) + \cdots + m_p(X_{ip})) \right)^2$$

subject to m_j smooth.

Additive models: Backfitting Algorithm

Input: Data (X_i, Y_i)

Iterate until convergence:

For each $j = 1, \dots, p$:

Compute residual: $R_j = Y - \sum_{k \neq j} \hat{m}_k(X_k)$

Smooth $\hat{m}_j = S_j R_j$

Output: Estimator $\hat{m}(X_i) = \sum_j \hat{m}_j(X_{ij})$.

Here, $S_j R$ is any 1-dimensional nonparametric regression smoother

But what if p is large?

Sparse Additive Models

Additive Model: $Y_i = \sum_{j=1}^p m_j(X_{ij}) + \varepsilon_i, \quad i = 1, \dots, n$

High dimensional: $n \ll p$, with most $m_j = 0$.

Optimization:

$$\begin{aligned} & \text{minimize} && \mathbb{E} \left(Y - \sum_j m_j(X_j) \right)^2 \\ & \text{subject to} && \sum_{j=1}^p \sqrt{\mathbb{E}(m_j^2)} \leq L_n \\ & && \mathbb{E}(m_j) = 0 \end{aligned}$$

This generalizes the lasso!

Sparse Backfitting Algorithm

Input: Data (X_i, Y_i) , regularization parameter λ .

Iterate until convergence:

For each $j = 1, \dots, p$:

Compute residual: $R_j = Y - \sum_{k \neq j} \hat{m}_k(X_k)$

Smooth $\hat{m}_j = \mathcal{S}_j R_j$

Estimate norm: $s_j = \sqrt{\mathbb{E}(\hat{m}_j^2)}$

Soft-threshold: $\hat{m}_j \leftarrow \left[1 - \frac{\lambda}{\hat{s}_j} \right]_+ \hat{m}_j$

Output: Estimator $\hat{m}(X_i) = \sum_j \hat{m}_j(X_{ij})$.

This generalizes coordinate descent algorithm from last time.

Kernel density estimation

To estimate a density, use the same idea behind kernel smoothing:

$$\begin{aligned}\hat{f}(x) &= \frac{1}{n} \sum_{i=1}^n K_h(X_i, x) \\ &= \frac{1}{n} \sum_{i=1}^n \frac{1}{h} K\left(\frac{X_i - x}{h}\right)\end{aligned}$$

We require that $\int K(u) du = 1$ and $K \geq 0$ is symmetric around zero (an even function).

This places a “bump” around each data point.

Kernel density estimation

In p dimensions:

$$\begin{aligned}\hat{f}(x) &= \frac{1}{n} \sum_{i=1}^n K_h(X_i, x) \\ &= \frac{1}{n h^p} \sum_{i=1}^n K\left(\frac{X_i - x}{h}\right)\end{aligned}$$

We require that $\int K(u) du = 1$ and K is symmetric around zero.

This places a “bump” around each data point.

Kernel density estimation

The bias-variance tradeoff:

$$\text{bias}^2(x) \approx h^4$$

$$\text{var}(x) \approx \frac{1}{n h^p}$$

Note that the variance scales according to the expected number of data points in a cube of side length h in p -dimensions.

We'll go through the calculation of this on the board.

Back to regression

Using a kernel density estimator, the “plug-in” regression estimate gives us back the kernel smoother:

$$\begin{aligned}\hat{m}(x) &= \int y \hat{f}(y | x) dy \\ &= \frac{\int y \hat{f}(x, y) dy}{\hat{f}(x)} \\ &= \frac{\sum_i Y_i K_h(X_i, x)}{\sum_i K_h(X_i, x)}\end{aligned}$$

Summary

- Smoothing methods compute local averages, weighting points by a kernel
- Shape of the kernel doesn't matter
- Mercer kernels using penalization rather than smoothing
- Defining property: Matrix \mathbb{K} is always positive-definite (non-negative eigenvalues)
- Equivalent to a type of ridge regression in function space
- The curse of dimensionality limits use of both approaches to low dimensions
- A compromise between nonparametric and linear models is to use additive models