

S&DS 365 / 665  
Intermediate Machine Learning

# Convolutional Neural Networks

February 14

Yale

# Topics for today

- Recap: Neural networks for classification
- Convolution operations
- Mechanics of convolutional networks
- Filters and pooling and flattening (oh my!)
- Example: Classifying Ca2+ brain scans
- Other examples

# Notes

- Animations from tutorial “Let’s Code Convolutional Neural Network in plain NumPy: Mysteries of Neural Networks” by Piotr Skalski
- [https://towardsdatascience.com/  
lets-code-convolutional-neural-network-in-plain-numpy-ce48e732f5d5](https://towardsdatascience.com/lets-code-convolutional-neural-network-in-plain-numpy-ce48e732f5d5)
- To see the animations, view this pdf using Adobe Acrobat

# Convolution

Mathematically, a *convolution* sweeps a function  $K(u)$  over the values of some function  $f$ :

$$K * f(x) = \int K(x - u) f(u) du$$

If  $K$  is a Gaussian bump, then this blurs or smooths out the values.

Note: Same operation we saw with smoothing kernels!

# Examples

convolve\_demo.ipynb gives examples



## Two principles

Two of the ideas behind convolutional neural networks (CNNs):

- ① Parameter tying: Let neurons share the same weights
- ② Kernel learning: The kernels should not be fixed, but rather *learned* from the data

CNNs are most commonly used for image data, but can be used whenever there is some kind of spatial structure to the data (e.g. DNA)

# Tensors

A tensor is just a multidimensional array\*. Like a matrix, but with more than just a pair of indices for row and columns.

In Python, a tensor is represented in `numpy` using the type `numpy.ndarray`

TensorFlow also has built-in types to represent tensors

---

\* In mathematics, there is a little more structure to tensors.

# Tensors

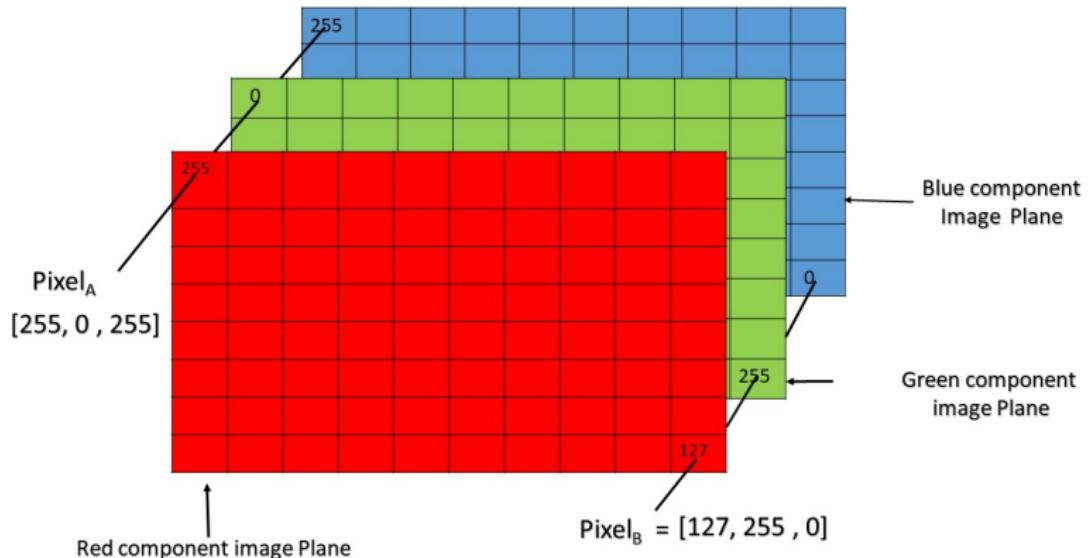
```
x = np.array(np.arange(24)).reshape(2, 3, 4)
print(x)
print(x[:, :, 0])

[[[ 0   1   2   3]
 [ 4   5   6   7]
 [ 8   9  10  11]]

[[12  13  14  15]
 [16  17  18  19]
 [20  21  22  23]]]

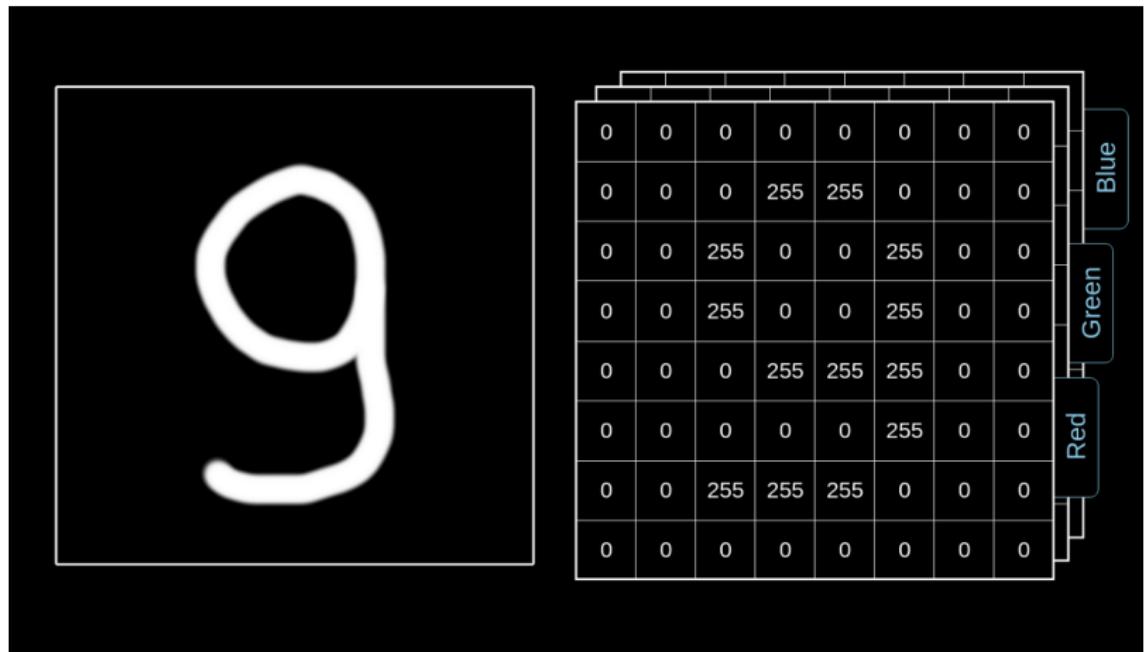
[[ 0   4   8]
 [12  16  20]]
```

# Images as tensors: Channels



Pixel of an RGB image are formed from the corresponding pixel of the three component images

# Images as tensors: Channels



# Convolutional layer

- In a convolutional layer, a set of *kernels* are “swept over” or convolved over the input
- The kernels are just weights that are trained
- The result of multiplying the kernel by the input tensor is called a *feature map*
- A convolutional layer is equivalent to a standard fully connected layer, but where the weights are “tied” (constrained to be equal) across different neurons

# Convolutional layer

# Adding nonlinearity

- So far everything is linear
- An activation function is used to add a nonlinearity
- Each filter also has an intercept (a single number), added before applying the activation function
- The same activation function is used for all of the kernels

# Pooling

- The point of a convolutional layer is to learn “features” of the input
- We don’t care too much exactly where they appear
- “Pooling” the feature map reduces the dimension and encodes the approximate location of features
- Two common types: max pooling and average pooling
- When using max pooling, to update weights in backpropagation, the locations of the features can be noted

# Pooling

# How features are built up

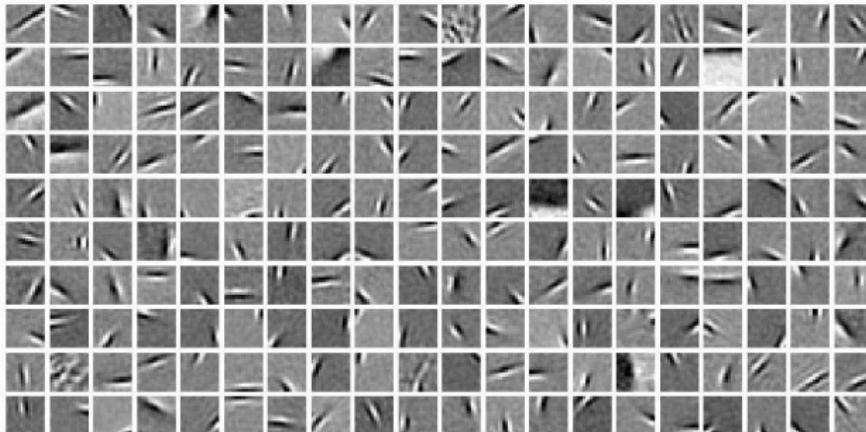
- Each kernel has a weight for each spatial position and channel
- Features from previous layers are recorded in the channels
- So the next layer can have a kernel that says

*“Oh, I see that feature 17 from the previous layer is active in position (1,1), feature 42 is in position (1,2), feature 5 is in position (2,1), and feature 11 is in position (2,2)”*

- These are compound features that are built up

# How features are built up

- Filters in the first layer may look like this:



- Filters in second layer can then capture more complex shapes

# Pooling and backpropagation

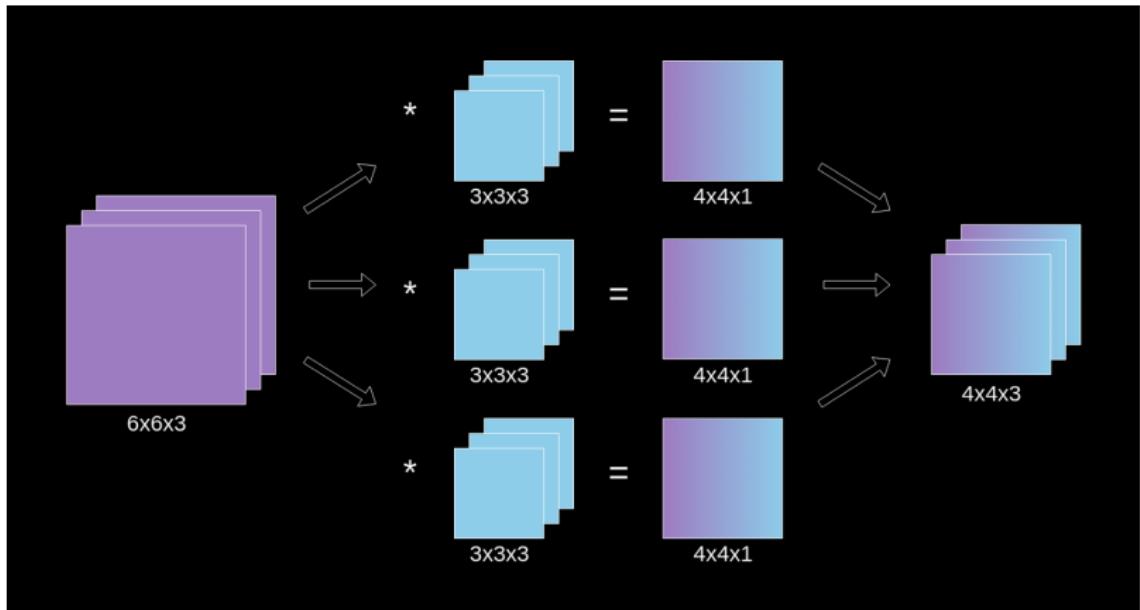
- A forward pass through the layers is used to make predictions, and store intermediate results to be used later
- A backward pass propagates the errors in prediction back through the layers, and computes changes to the weights for the gradient step

# **Forward and backward passes**

# Applying kernels across channels

- Typically each kernel (aka filter) has the same number of channels as the input tensor
- A kernel is applied at each location by multiplying component-wise and then summing
- The feature map for each kernel is then one channel in the output tensor
- So, the number of channels in the output is the number of kernels

# Applying kernels across channels



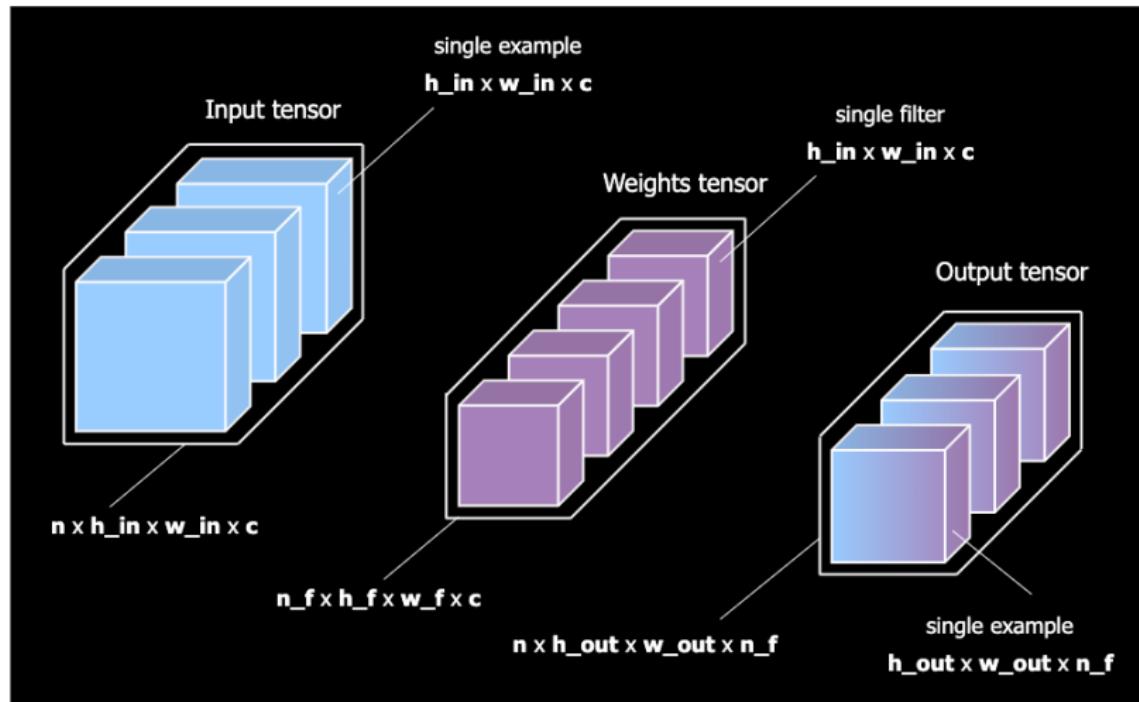
# Convolutional layer: Tensor shapes

- Data points are processed in “mini-batches” of  $n$  items
- The first dimension of the input and output tensors at each layer is always  $n$
- The last dimension of the output is the number of kernels (aka filters)

---

Can you find the typo in the next figure?

# Convolutional layer: Tensor shapes



Can you find the typo in this figure?

# Dropout

- “Dropout” is a type of regularization that is easy to incorporate
- Idea: Randomly zero out entries in a tensor
- Helps guard against overfitting
- Usual form:
  - ▶ Independently zero out each entry with probability  $\varepsilon$
  - ▶ Divide by  $(1 - \varepsilon)$

# Dropout

---

This illustrates something a bit different than the standard dropout

# Flattening

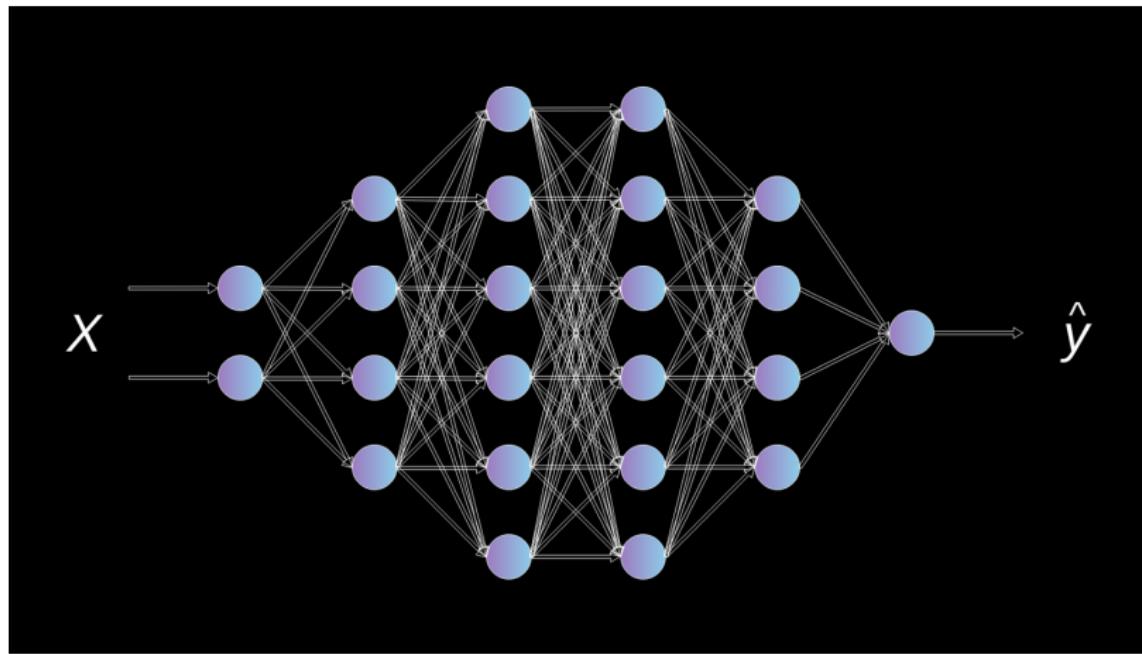
- After one or more convolutional layers, we usually add “dense” layers
- These are the standard types of layers in a neural network, with each output neuron connected to each input neuron, with a tuneable weight
- In CNNs, the convolutional layers build up a set of “features” and then these features are used to make a prediction

# Flattening

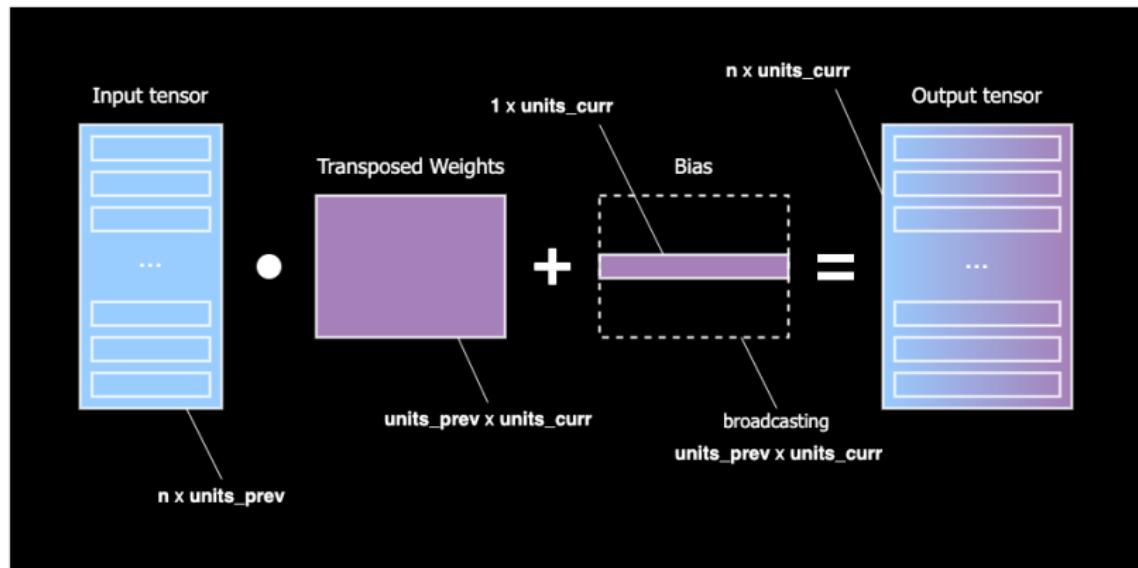
---

This is the same as the `numpy.flatten()` operation

# Dense layers



# Dense operations

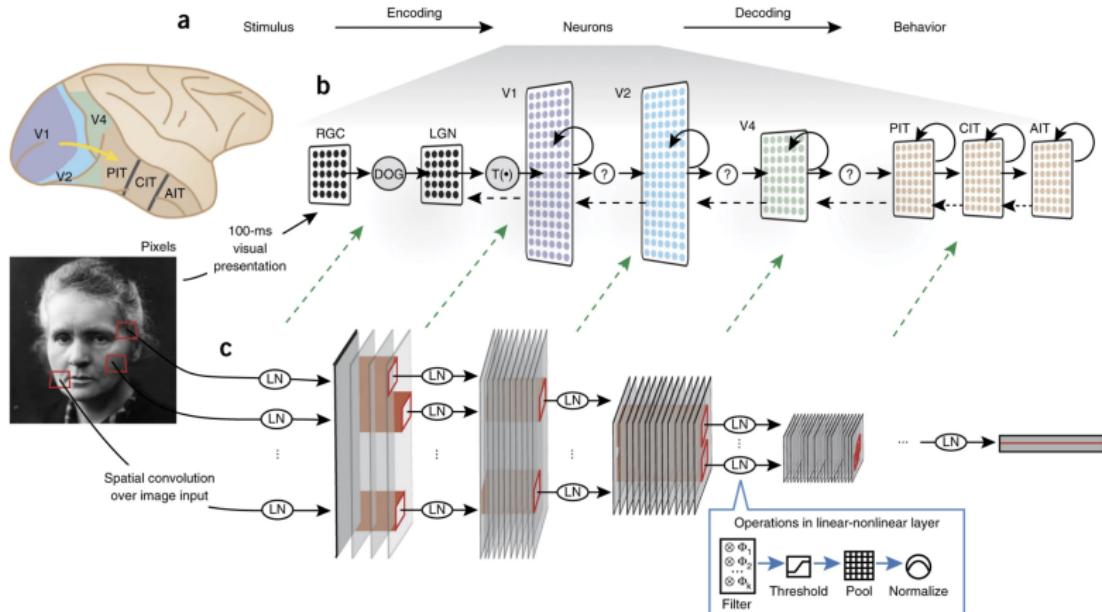


Again, this is not showing the use of an activation function

# CNNs: Biological analogies

- Arguments have been made that common architectures for CNNs used for image processing mirror the organization and function of visual cortex in many species

# CNNs: Biological analogies



Using goal-driven deep learning models to understand sensory cortex, Yamins and DiCarlo, 2016,  
<https://www.nature.com/articles/nn.4244>

## Problem 4: Brain food (20 points)

This problem gives you some experience with convolutional neural networks for image classification using [TensorFlow](#).

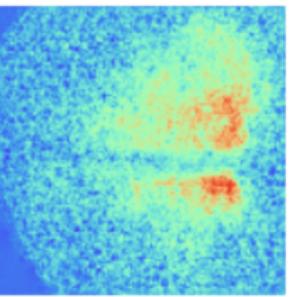
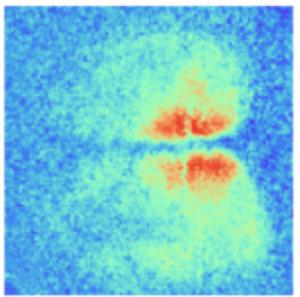
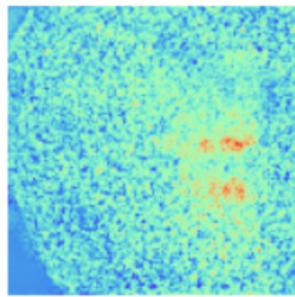
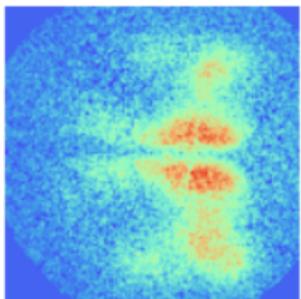
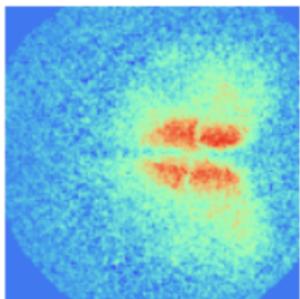
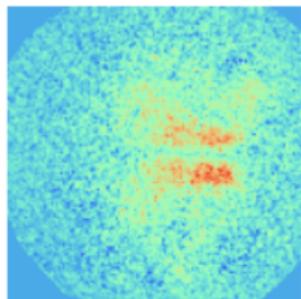
The classification task is to discriminate real optical images of brain activity in mice from fake images that were constructed using a [generative adversarial network \(GAN\)](#). A paper on the underlying imaging technologies developed by Yale researchers is [here](#).

For this problem we'll walk you through the following steps:

- Downloading the data
- Loading the data
- Displaying some sample images
- Building a classification model using a simple CNN

After this, your task will be to improve upon this baseline model by building, training, and evaluating two more CNNs.

# Task: Real or Fake?



# Task: Real or Fake?

- The real images are optical images of neural activity at the surface of the cortex in transgenic mice
- The fake images were generated by a “deconvolutional” neural network trained as part of a GAN
- The discriminator network is similar to the networks you’ll train

# Baseline model

```
model = models.Sequential()
model.add(layers.Conv2D(32, (5, 5), activation='relu', input_shape=(128, 128, 1)))
model.add(layers.MaxPooling2D((4, 4)))
model.add(layers.Flatten())
model.add(layers.Dense(2))
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d (Conv2D)	(None, 124, 124, 32)	832
<hr/>		
max_pooling2d (MaxPooling2D)	(None, 31, 31, 32)	0
<hr/>		
flatten (Flatten)	(None, 30752)	0
<hr/>		
dense (Dense)	(None, 2)	61506
<hr/>		

Total params: 62,338

Trainable params: 62,338

Non-trainable params: 0

# Baseline model

First layer:

- Input:  $128 \times 128$  images, 1 channel
- First layer: 32 filters, each of size  $5 \times 5$
- Output of first layer: tensor of shape  $124 \times 124 \times 32$
- Number of parameters:  $(5 \times 5 + 1) \times 32 = 832$
- Activation function: ReLU
- Output of first layer: tensor of shape  $124 \times 124 \times 32$

# Baseline model

Second layer:

- Maximum of  $4 \times 4$  regions in feature map
- Output: Tensor of shape  $31 \times 31 \times 32$  since  $124/4 = 31$
- Number of parameters: None!

# Baseline model

Third layer:

- Flattened version of previous layer
- Number of neurons:  $31 \times 31 \times 32 = 30,752$
- Number of parameters: None!

# Baseline model

Final layer:

- Dense connections
- Two neurons in output (“fake” and “real”)
- Number of parameters:  $30,752 \times 2 = 61,506$

# Baseline model

Overall:

- Number of parameters:  $832 + 61,506 = 62,338$
- Most come from dense layer (98.7%)

# Let's build another model

```
model = models.Sequential()
model.add(layers.Conv2D(32, (5, 5), activation='relu', input_shape=(128, 128, 1)))
model.add(layers.MaxPooling2D((4, 4)))
model.add(layers.Flatten())
model.add(layers.Dense(2))
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 124, 124, 32)	832
max_pooling2d (MaxPooling2D)	(None, 31, 31, 32)	0
flatten (Flatten)	(None, 30752)	0
dense (Dense)	(None, 2)	61506

Total params: 62,338  
Trainable params: 62,338  
Non-trainable params: 0

# Summary

- A convolutional layer applies a set of kernels (filters) across the input
- Each kernel is a tensor of smaller dimension than the input
- The values of the kernels are learned by backpropagation
- The convolutional layers define a set of features that are then used to make predictions in one or more dense layers
- CNNs are attractive whenever the data have “spatial” structure where the same features might appear at many different locations